

Tractor Sales Forecasting

Himank Jain

17/02/2020

Buiseness Problem:

PowerHorse, a tractor and farm equipment manufacturing company, was established a few years after World War II. The company has shown a consistent growth in its revenue from tractor sales since its inception. However, over the years the company has struggled to keep it's inventory and production cost down because of variability in sales and tractor demand. The management at PowerHorse is under enormous pressure from the shareholders and board to reduce the production cost.

Develop models to forecast next 2 years sales.

Data:

The dataset consists of 144 observations having the total monthwise sales data of Tractors for a period of past 12 years.

Loading Required Libraries:

```
library(data.table)
library(ggplot2)
library(fpp2)
library(forecast)
library(stats)
library(tseries)
```

Exploratory Analysis:


```
sales<-read.csv("Tractor-Sales.csv")
head(sales)
```

```
##   Month.Year Number.of.Tractor.Sold
## 1   Jan-03             141
## 2   Feb-03             157
## 3   Mar-03             185
## 4   Apr-03             199
## 5   May-03             203
## 6   Jun-03             189
```

```
# Converting data to timeseries:
sales_ts=ts(sales[,2],start=c(2003,1),frequency=12)
sales_ts
```

```
##      Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
## 2003 141 157 185 199 203 189 207 207 171 150 138 165
## 2004 145 168 197 208 210 209 238 238 199 168 152 196
## 2005 183 200 249 251 289 249 279 279 232 204 194 232
## 2006 215 239 270 279 307 305 322 339 263 241 229 272
## 2007 247 261 330 362 385 340 370 381 299 266 239 281
## 2008 257 250 329 350 393 370 423 410 326 289 270 321
## 2009 305 310 374 414 454 441 510 486 393 345 315 389
## 2010 358 368 444 482 534 524 578 567 447 386 360 428
## 2011 397 400 498 536 596 591 651 654 509 437 406 470
## 2012 428 423 507 536 610 609 687 707 509 452 412 472
## 2013 454 455 568 610 706 661 767 783 583 513 481 567
## 2014 525 520 587 710 793 749 871 848 640 581 519 605
```


```
plot(sales_ts,xlab="Years",ylab="Tractor Sales",main="Tractor Sales Data")
```



Sesonalty:


```
ggseasonplot(sales_ts,year.labels = TRUE,year.labels.left = TRUE)+ylab("Degree")+
  ggtitle("Season Plot: Tractor Sales")
```

Season Plot: Tractor Sales



```
ggseasonplot(sales_ts,polar=TRUE )+ylab("Degree")+
  ggtitle("Season Plot: Tractor Sales")
```


Season Plot: Tractor Sales



be the peak months for sales.

July & August seem to

```
monthplot(sales_ts)
```



lar components at April 2008 and Feb 2008

Decomposition:

```
sales_decompose<-decompose(sales_ts,type='multiplicative')
sales_decompose
```

```
## $x
##      Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
## 2003 141 157 185 199 203 189 207 207 171 150 138 165
## 2004 145 168 197 208 210 209 238 238 199 168 152 196
## 2005 183 200 249 251 289 249 279 279 232 204 194 232
## 2006 215 239 270 279 307 305 322 339 263 241 229 272
## 2007 247 261 330 362 385 340 370 381 299 266 239 281
## 2008 257 250 329 350 393 370 423 410 326 289 270 321
## 2009 305 310 374 414 454 441 510 486 393 345 315 389
## 2010 358 368 444 482 534 524 578 567 447 386 360 428
## 2011 397 400 498 536 596 591 651 654 509 437 406 470
## 2012 428 423 507 536 610 609 687 707 509 452 412 472
## 2013 454 455 568 610 706 661 767 783 583 513 481 567
## 2014 525 520 587 710 793 749 871 848 640 581 519 605
##
## $seasonal
##      Jan      Feb      Mar      Apr      May      Jun      Jul
## 2003 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
## 2004 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
## 2005 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
## 2006 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
## 2007 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
## 2008 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
## 2009 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
## 2010 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
## 2011 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
## 2012 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
## 2013 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
## 2014 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238
##      Aug      Sep      Oct      Nov      Dec
## 2003 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
## 2004 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
## 2005 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
## 2006 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
## 2007 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
## 2008 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
## 2009 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
## 2010 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
## 2011 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
## 2012 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
## 2013 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
## 2014 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
##
## $trend
##      Jan      Feb      Mar      Apr      May      Jun      Jul      Aug
## 2003      NA      NA      NA      NA      NA      NA 176.1667 176.7917
## 2004 182.5417 185.1250 187.5833 189.5000 190.8333 192.7083 195.5833 198.5000
```

```

## 2005 219.3750 222.7917 225.8750 228.7500 232.0000 235.2500 238.0833 241.0417
## 2006 254.7083 259.0000 262.7917 265.6250 268.6250 271.7500 274.7500 277.0000
## 2007 301.2500 305.0000 308.2500 310.7917 312.2500 313.0417 313.8333 313.7917
## 2008 317.6250 321.0417 323.3750 325.4583 327.7083 330.6667 334.3333 338.8333
## 2009 365.0417 371.8333 377.7917 382.9167 387.1250 391.8333 396.8750 401.5000
## 2010 431.8333 438.0417 443.6667 447.6250 451.2083 454.7083 457.9583 460.9167
## 2011 485.0417 491.7083 497.9167 502.6250 506.6667 510.3333 513.3750 515.6250
## 2012 521.5000 525.2083 527.4167 528.0417 528.9167 529.2500 530.4167 532.8333
## 2013 561.0833 567.5833 573.8333 579.4583 584.8750 591.7083 598.6250 604.2917
## 2014 635.8333 642.8750 647.9583 653.1667 657.5833 660.7500      NA      NA
##      Sep      Oct      Nov      Dec
## 2003 177.7500 178.6250 179.2917 180.4167
## 2004 202.0000 205.9583 211.0417 216.0000
## 2005 243.5417 245.5833 247.5000 250.5833
## 2006 280.4167 286.3750 293.0833 297.7917
## 2007 313.2917 312.7500 312.5833 314.1667
## 2008 343.2083 347.7500 352.9583 358.4583
## 2009 406.8333 412.5833 418.7500 425.5417
## 2010 464.5000 469.0000 473.8333 479.2083
## 2011 516.9583 517.3333 517.9167 519.2500
## 2012 536.7083 542.3333 549.4167 555.5833
## 2013 607.7917 612.7500 620.5417 627.8333
## 2014      NA      NA      NA      NA
##
## $random
##      Jan      Feb      Mar      Apr      May      Jun      Jul
## 2003      NA      NA      NA      NA      NA      NA 0.9506481
## 2004 0.9647844 1.0754101 1.0373675 1.0157964 0.9280175 0.9675514 0.9845058
## 2005 1.0131838 1.0638023 1.0889096 1.0154655 1.0505119 0.9442738 0.9480874
## 2006 1.0252260 1.0935233 1.0148756 0.9720482 0.9637914 1.0012869 0.9481812
## 2007 0.9958506 1.0140761 1.0574784 1.0779330 1.0397990 0.9689581 0.9538407
## 2008 0.9827493 0.9228019 1.0049631 0.9952340 1.0113376 0.9982505 1.0236085
## 2009 1.0148029 0.9879689 0.9778672 1.0005726 0.9889987 1.0040734 1.0396559
## 2010 1.0069112 0.9955482 0.9885231 0.9965182 0.9980568 1.0280793 1.0211159
## 2011 0.9941130 0.9640117 0.9879463 0.9869002 0.9920077 1.0331462 1.0259341
## 2012 0.9968131 0.9544180 0.9495433 0.9393970 0.9725987 1.0265608 1.0478829
## 2013 0.9827720 0.9499738 0.9777396 0.9742270 1.0179648 0.9966026 1.0366059
## 2014 1.0028600 0.9585322 0.8948529 1.0059743 1.0169825 1.0112834      NA
##      Aug      Sep      Oct      Nov      Dec
## 2003 0.9518221 1.0000639 1.0039910 1.0054376 1.0126675
## 2004 0.9746837 1.0241013 0.9752383 0.9408307 1.0047590
## 2005 0.9409344 0.9902753 0.9931437 1.0239120 1.0251688
## 2006 0.9948719 0.9749743 1.0061501 1.0206586 1.0113837
## 2007 0.9870311 0.9921186 1.0168693 0.9987762 0.9903890
## 2008 0.9836599 0.9874179 0.9935999 0.9992553 0.9915759
## 2009 0.9840065 1.0041931 0.9997427 0.9826341 1.0122018
## 2010 1.0000182 1.0003755 0.9840004 0.9924600 0.9889606
## 2011 1.0310768 1.0235368 1.0099312 1.0240054 1.0022613
## 2012 1.0786369 0.9858723 0.9964441 0.9795610 0.9407027
## 2013 1.0533250 0.9971375 1.0009557 1.0125355 0.9999961
## 2014      NA      NA      NA      NA      NA
##
## $figure
## [1] 0.8233333 0.8438594 1.0123702 1.0805564 1.1857930 1.1209126 1.2360238

```

```
## [8] 1.2301349 0.9619639 0.8364099 0.7655329 0.9031095
##
## $type
## [1] "multiplicative"
##
## attr(,"class")
## [1] "decomposed.ts"
```

Decomposition gives the following components:


\$x: The original data

\$seasonal: Sales percentage relative to annual trend. Therefore, In January sales are 18% less than annual trend.

\$trend: continuous trend of sales yearly


\$random: Peculiarity i.e. Variation in data having accounted for seasonality & trend.

```
plot(sales_decompose)
```



- The trend is generally growing over the years.
- The random error is within 10% of total sales for the years.
- Therefore, If the company has a high budget, we can confidently plan the production to be the forecasted sales $\pm 10\%$.
- If costs are high and company has a low budget, then we can plan the production to be 450 and be confident all tractors are going to be sold.


```
## additive model on logarithm scale.
sales_decom_log<-stl(log10(sales_ts),s.window='p')
plot(sales_decom_log)
```



Train Test Split:

```
ts_train<-window(sales_ts,start=c(2003,1),end=c(2012,12),freq=12)
ts_test<-window(sales_ts,start=c(2013,1),freq=12)
```

```
autoplot(ts_train,,series="Train")+autolayer(ts_test,series="Test")+ggtitle("Tractor sales train & test")
```

Forecasting:

Random Walk with Drift:

A random walk is defined as a process where the current value of a variable is composed of the past value plus an error term defined as a white noise (a normal variable with zero mean and variance one). Algebraically a random walk is represented as follows:

$$y_t = y_{t-1} + \epsilon_t$$


```
ts_decompose_train_log<-stl(log10(ts_train),s.window='p')
ts_train_stl<-forecast(ts_decompose_train_log,method='rwdrift',h=24)
ts_train_stl
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2013	2.634223	2.614268	2.654178	2.603705	2.664742
## Feb 2013	2.656170	2.627831	2.684509	2.612830	2.699511
## Mar 2013	2.740345	2.705493	2.775197	2.687043	2.793647
## Apr 2013	2.768221	2.727812	2.808631	2.706420	2.830023
## May 2013	2.806446	2.761082	2.851811	2.737068	2.875825
## Jun 2013	2.786750	2.736855	2.836646	2.710441	2.863059
## Jul 2013	2.834038	2.779928	2.888149	2.751284	2.916793
## Aug 2013	2.835597	2.777520	2.893674	2.746775	2.924419
## Sep 2013	2.736191	2.674346	2.798035	2.641608	2.830773
## Oct 2013	2.679677	2.614231	2.745123	2.579586	2.779768

## Nov 2013	2.645540	2.576632	2.714448	2.540154	2.750925
## Dec 2013	2.722442	2.650192	2.794693	2.611945	2.832940
## Jan 2014	2.682723	2.607234	2.758213	2.567273	2.798174
## Feb 2014	2.704671	2.626034	2.783308	2.584406	2.824936
## Mar 2014	2.788845	2.707140	2.870550	2.663888	2.913802
## Apr 2014	2.816722	2.732021	2.901423	2.687183	2.946261
## May 2014	2.854947	2.767314	2.942580	2.720923	2.988970
## Jun 2014	2.835251	2.744744	2.925758	2.696832	2.973669
## Jul 2014	2.882539	2.789210	2.975867	2.739805	3.025272
## Aug 2014	2.884097	2.787996	2.980199	2.737122	3.031072
## Sep 2014	2.784691	2.685860	2.883522	2.633542	2.935840
## Oct 2014	2.728177	2.626657	2.829697	2.572916	2.883439
## Nov 2014	2.694040	2.589868	2.798212	2.534723	2.853357
## Dec 2014	2.770943	2.664154	2.877732	2.607623	2.934263


```
plot(ts_train_stl)
```

Forecasts from STL + Random walk with drift



```
vec2<-10^(cbind(log10(ts_test),as.data.frame(ts_train_stl)[,1]))
ts.plot(vec2,col=c('blue','red'),main='Tractor sales Actual vs Forecast')
```

Tractor sales Actual vs Forecast



```
RMSE<-round(sqrt(sum(((vec2[,1]-vec2[,2])^2)/length(vec2[,1]))),4)
MAPE<-round(mean(abs(vec2[,1]-vec2[,2])/vec2[,1]),4)
paste("Accuracy Measures: RMSE: ",RMSE,"and MAPE: ",MAPE)
```

```
## [1] "Accuracy Measures: RMSE: 53.5697 and MAPE: 0.0687"
```

Therefore The forecast on average is about 6.87% away from actual sales. Underpredicting by 6.87%

```
10~as.data.frame(ts_train_stl)
```

```
##          Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95
## Jan 2013    430.7478 411.4036 451.0015 401.5176 462.1059
## Feb 2013    453.0753 424.4548 483.6255 410.0434 500.6231
## Mar 2013    549.9774 507.5662 595.9325 486.4554 621.7943
## Apr 2013    586.4370 534.3324 643.6224 508.6510 676.1185
## May 2013    640.3927 576.8754 710.9035 545.8428 751.3203
## Jun 2013    611.9986 545.5753 686.5088 513.3830 729.5572
## Jul 2013    682.3988 602.4595 772.9452 564.0057 825.6444
## Aug 2013    684.8525 599.1280 782.8426 558.1813 840.2698
## Sep 2013    544.7418 472.4397 628.1091 438.1350 677.2882
## Oct 2013    478.2741 411.3685 556.0613 379.8272 602.2372
## Nov 2013    442.1196 377.2524 518.1405 346.8602 563.5404
## Dec 2013    527.7672 446.8813 623.2935 409.2090 680.6747
## Jan 2014    481.6410 404.7943 573.0765 369.2095 628.3102
## Feb 2014    506.6065 422.7013 607.1667 384.0657 668.2453
## Mar 2014    614.9577 509.4955 742.2500 461.1990 819.9780
## Apr 2014    655.7250 539.5362 796.9350 486.6117 883.6107
## May 2014    716.0556 585.2125 876.1528 525.9246 974.9224
```

## Jun 2014	684.3067	555.5763	842.8648	497.5447	941.1730
## Jul 2014	763.0248	615.4747	945.9477	549.2944	1059.9178
## Aug 2014	765.7684	613.7558	955.4307	545.9117	1074.1685
## Sep 2014	609.1035	485.1320	764.7549	430.0727	862.6614
## Oct 2014	534.7825	423.3086	675.6120	374.0378	764.6081
## Nov 2014	494.3564	388.9272	628.3650	342.5493	713.4396
## Dec 2014	590.1233	461.4809	754.6261	405.1567	859.5330

Therefore 95% confidence interval : forecast +/- 1.96 * RMSE i.e. For January 2013 Forecast interval: 430 +/- 1.96 * 53.56

Exponential Smoothing:

Exponential smoothing is a popular forecasting method for short-term predictions. Such forecasts of future values are based on past data whereby the most recent observations are weighted more than less recent observations. As part of this weighting, constants are being smoothed. This is different from the simple moving average method, in which every data point has equal weight in the average calculation. Exponential smoothing introduces the idea of building a forecasted value as the average figure from differently weighted data points for the average calculation.

There are different exponential smoothing methods that differ from each other in the components of the time series that are modeled. * Single Exponential Smoothing * Double Exponential Smoothing * Triple Exponential Smoothing

Simple exponential smoothing (SES) uses only one smoothing constant, double exponential smoothing or Holt exponential smoothing uses two smoothing constants and triple exponential smoothing or Holt-Winters exponential smoothing accordingly uses three smoothing constants.

Single Exponential Smoothing:

Simple exponential smoothing assumes that the time series data has only a level and some error (or remainder) but no trend or seasonality The smoothing parameter α determines the distribution of weights of past observations and with that how heavily a given time period is factored into the forecasted value. If the smoothing parameter is close to 1, recent observations carry more weight and if the smoothing parameter is closer to 0

$$F_t = F_{t-1} + \alpha * (A_{t-1} - F_{t-1})$$

or

$$F_t = \alpha * A_{t-1} + (1 - \alpha) * F_{t-1}$$

F_{t-1} = forecast for the previous period,

A_{t-1} = Actual demand for the period,

α = weight (between 0 and 1). The closer to zero, the smaller the weight.


```
ts_train_ses<-ses(ts_train,hs=24)
ts_train_ses
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2013	471.994	411.0351	532.9529	378.7654	565.2226
## Feb 2013	471.994	385.7894	558.1986	340.1555	603.8325
## Mar 2013	471.994	366.4172	577.5708	310.5282	633.4598

```
## Apr 2013      471.994 350.0854 593.9026 285.5509 658.4371
## May 2013      471.994 335.6967 608.2913 263.5453 680.4427
## Jun 2013      471.994 322.6883 621.2997 243.6506 700.3374
## Jul 2013      471.994 310.7258 633.2622 225.3556 718.6324
## Aug 2013      471.994 299.5914 644.3966 208.3269 735.6611
## Sep 2013      471.994 289.1336 654.8544 192.3332 751.6548
## Oct 2013      471.994 279.2425 664.7455 177.2060 766.7820
```

```
plot(ts_train_ses)
```

Forecasts from Simple exponential smoothing



Observed that all Point forecasts are the same.

Holt Method (Double Exponential Smoothing):

Holt exponential smoothing is a time series forecasting approach that fits time series data with an overall level as well as a trend. Additionally, to simple exponential smoothing, which uses smoothing parameter α only there is also a β smoothing parameter for the exponential decay of the modeled trend component. This β smoothing parameter ranges between 0 and 1, with higher values indicating more weight to recent observations. It accounts for level and trend but not seasonality.


$$b_t = \gamma(S_t - S_{t-1}) + (1 - \gamma)b_{t-1}$$

```
ts_train_holt<-holt(ts_train,h=24)
ts_train_holt
```

```
##          Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95
## Jan 2013      475.2876 413.8077 536.7674 381.26220 569.3129
## Feb 2013      478.5814 391.6360 565.5269 345.60982 611.5530
## Mar 2013      481.8753 375.3856 588.3650 319.01336 644.7372
## Apr 2013      485.1691 362.2001 608.1382 297.10414 673.2342
```

## May 2013	488.4630	350.9728	625.9532	278.18993	698.7361
## Jun 2013	491.7569	341.1365	642.3773	261.40286	722.1109
## Jul 2013	495.0507	332.3539	657.7476	246.22739	743.8741
## Aug 2013	498.3446	324.4055	672.2838	232.32764	764.3616
## Sep 2013	501.6385	317.1386	686.1383	219.47030	783.8066
## Oct 2013	504.9323	310.4423	699.4224	207.48552	802.3792
## Nov 2013	508.2262	304.2327	712.2197	196.24507	820.2073
## Dec 2013	511.5201	298.4447	724.5955	185.64935	837.3908
## Jan 2014	514.8139	293.0264	736.6015	175.61914	854.0087
## Feb 2014	518.1078	287.9359	748.2797	166.09023	870.1254
## Mar 2014	521.4017	283.1386	759.6648	157.00968	885.7937
## Apr 2014	524.6955	278.6055	770.7856	148.33327	901.0578
## May 2014	527.9894	274.3122	781.6666	140.02359	915.9552
## Jun 2014	531.2833	270.2378	792.3287	132.04866	930.5179
## Jul 2014	534.5771	266.3642	802.7900	124.38089	944.7734
## Aug 2014	537.8710	262.6758	813.0662	116.99628	958.7457
## Sep 2014	541.1649	259.1588	823.1709	109.87380	972.4559
## Oct 2014	544.4587	255.8010	833.1164	102.99489	985.9226
## Nov 2014	547.7526	252.5918	842.9134	96.34309	999.1621
## Dec 2014	551.0465	249.5214	852.5715	89.90369	1012.1892

```
plot(ts_train_holt)
```



Holt Winter's Method (triple exponential model):

Holt-Winters exponential smoothing is a time series forecasting approach that takes the overall level, trend and seasonality of the underlying dataset into account for its forecast. Hence, the Holt-Winters model has three smoothing parameters indicating the exponential decay from most recent to older observations: α for the level component, β for the trend component, and γ for the seasonality component. It accounts for level and trend and seasonality.

level:

$$L_t = \alpha * (y_t - S_{t-s}) + (1 - \alpha)(L_{t-1} + b_{t-1})$$

trend:

$$b_t = \beta * (L_t - L_{t-1}) + (1 - \beta) * b_{t-1}$$

seasonal:

$$S_t = \gamma(y_t - L_t) + (1 - \gamma)S_{t-s}$$

forecast:


$$F_{t+k} = L_t + k * b_t + S_{t+k-s}$$

```
ts_train_hw<-hw(ts_train,h=24,seasonal="multiplicative")
ts_train_hw
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2013	440.8777	416.6841	465.0713	403.8768	477.8786
## Feb 2013	442.2152	416.4192	468.0111	402.7637	481.6667
## Mar 2013	540.7288	507.3398	574.1179	489.6647	591.7929
## Apr 2013	582.7737	544.8184	620.7291	524.7261	640.8214
## May 2013	666.0299	620.4189	711.6408	596.2740	735.7858
## Jun 2013	663.9650	616.2837	711.6463	591.0428	736.8872
## Jul 2013	741.2300	685.5440	796.9160	656.0656	826.3944
## Aug 2013	747.1779	688.5796	805.7761	657.5596	836.7962
## Sep 2013	543.3566	498.9564	587.7568	475.4523	611.2609
## Oct 2013	479.4786	438.7255	520.2318	417.1520	541.8052
## Nov 2013	437.5753	398.9513	476.1993	378.5050	496.6457
## Dec 2013	504.9576	458.7353	551.1799	434.2666	575.6486
## Jan 2014	471.6816	419.9860	523.3772	392.6200	550.7432
## Feb 2014	472.9365	419.7595	526.1135	391.6093	554.2637
## Mar 2014	578.0813	511.4267	644.7359	476.1419	680.0207
## Apr 2014	622.8040	549.1967	696.4113	510.2314	735.3766
## May 2014	711.5228	625.3641	797.6816	579.7544	843.2913
## Jun 2014	709.0645	621.1309	796.9980	574.5817	843.5472
## Jul 2014	791.2990	690.8398	891.7582	637.6599	944.9381
## Aug 2014	797.3708	693.7814	900.9603	638.9444	955.7973
## Sep 2014	579.6577	502.6271	656.6884	461.8495	697.4660
## Oct 2014	511.3378	441.8559	580.8197	405.0744	617.6011
## Nov 2014	466.4928	401.7011	531.2845	367.4024	565.5832
## Dec 2014	538.1485	461.7771	614.5199	421.3485	654.9484

```
plot(ts_train_hw)
```


Forecasts from Holt–Winters' multiplicative method



Holt Winter's method:

```
vec<-cbind(ts_test,as.data.frame(ts_train_hw)[,1])
ts.plot(vec,col=c("blue","red"),main="Tractor Sales: Actual vs Forecast")
```

Tractor Sales: Actual vs Forecast



Time

Underprediction! This means the trend seems to be increasing faster than what historical data is suggesting.


```
RMSE<-round(sqrt(sum(((vec[,1]-vec[,2])^2)/length(vec[,1]))),4)
MAPE<-round(mean(abs(vec[,1]-vec[,2])/vec[,1]),4)
paste("Accuracy Measures: RMSE: ",RMSE,"and MAPE: ",MAPE)
```

```
## [1] "Accuracy Measures: RMSE: 49.7553 and MAPE: 0.0703"
```


In analytics perspective, RMSE tracks variance while MAPE tracks bias.

Comparing rwdrift model to hold winter model, we observe that rwdrift model has less bias but high variance while holt winter model has less variance and high bias.

Stationarity:

Original data:

```
plot(sales_ts, xlab='Years', ylab = 'Tractor Sales')
```



Checking Stationarity using Augmented Dickey-Fuller test:

```
adf.test(sales_ts)
```

```
## Warning in adf.test(sales_ts): p-value smaller than printed p-value
```


```
##
## Augmented Dickey-Fuller Test
##
## data: sales_ts
## Dickey-Fuller = -12.783, Lag order = 5, p-value = 0.01
## alternative hypothesis: stationary
```

Therefore Data is not stationary. Let's make it stationary for ARIMA modelling.

Trend Removal:


Removing the upward trend through 1st order differencing.

```
plot(diff(sales_ts),ylab='Differenced Tractor Sales')
```




Log transform data to make data stationary on variance Now, the above series is not stationary on variance i.e. variation in the plot is increasing as we move towards the right of the chart. We need to make the series stationary on variance to produce reliable forecasts through ARIMA models.

```
plot(log10(sales_ts),ylab='Log (Tractor Sales)')
```



Difference log transform data to make data stationary on both mean and variance:

```
plot(diff(log10(sales_ts)),ylab='Differenced Log (Tractor Sales)')
```




The time series looks stationary. The Integrated part (I) of ARIMA model will be equal to 1 as we used 1st order difference to make series stationary.

ACF & PACF:


Plot ACF(Autocorrelation factor) and PACF(Partial Autocorrelation factor) to identify potential AR and MA components in the residuals.

```
par(mfrow=c(1,2))
acf(ts(diff(log10(sales_ts))),main='ACF Tractor Sales')
pacf(ts(diff(log10(sales_ts))),main='PACF Tractor Sales')
```

ACF Tractor Sales



PACF Tractor Sales



There are spikes beyond significant zones. Therefore residuals are not random. Hence AR MA models can be used to extract this information.

ARIMA:

Testing arima model on ts_test data:

```
ARIMAfit = auto.arima(log10(ts_train), approximation=FALSE,trace=FALSE)
summary(ARIMAfit)
```


```
## Series: log10(ts_train)
## ARIMA(0,1,1)(0,1,1)[12]
##
## Coefficients:
##          ma1      sma1
##       -0.3471  -0.5355
## s.e.   0.0996   0.0876
##
## sigma^2 estimated as 0.0002688:  log likelihood=287.06
## AIC=-568.12  AICc=-567.89  BIC=-560.11
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE
## Training set -6.71787e-05 0.01533676 0.01139577 -0.002628277 0.4559706
##              MASE      ACF1
## Training set 0.2131231 0.002369841
```

```
pred<-forecast(ARIMAfit,h=24)
vec<-10^cbind(log10(ts_test),as.data.frame(pred)[,1])
vec
```

```
##          log10(ts_test) as.data.frame(pred)[, 1]
## Jan 2013           454           439.3281
## Feb 2013           455           441.0230
## Mar 2013           568           536.5490
## Apr 2013           610           573.9486
## May 2013           706           643.5105
## Jun 2013           661           634.8851
## Jul 2013           767           710.3604
## Aug 2013           783           716.3897
## Sep 2013           583           541.0908
## Oct 2013           513           474.3412
## Nov 2013           481           436.3978
## Dec 2013           567           508.0340
## Jan 2014           525           469.7022
## Feb 2014           520           471.5143
## Mar 2014           587           573.6447
## Apr 2014           710           613.6301
## May 2014           793           688.0014
## Jun 2014           749           678.7796
## Jul 2014           871           759.4731
## Aug 2014           848           765.9192
## Sep 2014           640           578.5006
## Oct 2014           581           507.1361
## Nov 2014           519           466.5693
## Dec 2014           605           543.1583
```

```
ts.plot(vec,col=c("blue","red"),main="Tractor Sales: Actual vs Forecast")
```

Tractor Sales: Actual vs Forecast



```
RMSE<-round(sqrt(sum(((vec[,1]-vec[,2])^2)/length(vec[,1]))),4)
MAPE<-round(mean(abs(vec[,1]-vec[,2])/vec[,1]),4)
paste("Accuracy Measures: RMSE: ",RMSE,"and MAPE: ",MAPE)
```

```
## [1] "Accuracy Measures: RMSE: 61.0551 and MAPE: 0.0851"
```

Using arima to forecast values of 2015, 2016 & 2017

```
ARIMAfit = auto.arima(log10(sales_ts), approximation=FALSE,trace=FALSE)
summary(ARIMAfit)
```


```
## Series: log10(sales_ts)
## ARIMA(0,1,1)(0,1,1)[12]
##
## Coefficients:
##          ma1      sma1
##       -0.4047 -0.5529
## s.e.    0.0885  0.0734
##
## sigma^2 estimated as 0.0002571: log likelihood=354.4
## AIC=-702.79 AICc=-702.6 BIC=-694.17
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.0002410698 0.01517695 0.01135312 0.008335713 0.4462212 0.2158968
##              ACF1
## Training set 0.01062604
```

Forecasting:

```
par(mfrow = c(1,1))
pred = predict(ARIMAfit, n.ahead = 36) #36 months makes 3 years
pred
```

```
## $pred
##           Jan      Feb      Mar      Apr      May      Jun      Jul      Aug
## 2015 2.754168 2.753182 2.826608 2.880192 2.932447 2.912372 2.972538 2.970585
## 2016 2.796051 2.795065 2.868491 2.922075 2.974330 2.954255 3.014421 3.012468
## 2017 2.837934 2.836948 2.910374 2.963958 3.016213 2.996138 3.056304 3.054351
##           Sep      Oct      Nov      Dec
## 2015 2.847264 2.797259 2.757395 2.825125
## 2016 2.889147 2.839142 2.799278 2.867008
## 2017 2.931030 2.881025 2.841161 2.908891
##
## $se
##           Jan      Feb      Mar      Apr      May      Jun
## 2015 0.01603508 0.01866159 0.02096153 0.02303295 0.02493287 0.02669792
## 2016 0.03923008 0.04159145 0.04382576 0.04595157 0.04798329 0.04993241
## 2017 0.06386474 0.06637555 0.06879478 0.07113179 0.07339441 0.07558934
##           Jul      Aug      Sep      Oct      Nov      Dec
## 2015 0.02835330 0.02991723 0.03140337 0.03282229 0.03418236 0.03549035
## 2016 0.05180825 0.05361850 0.05536960 0.05706700 0.05871534 0.06031866
## 2017 0.07772231 0.07979828 0.08182160 0.08379608 0.08572510 0.08761165
```

```
plot(sales_ts,type='l',xlim=c(2004,2018),ylim=c(1,1600),xlab = 'Year',ylab = 'Tractor Sales')
lines(10^(pred$pred),col='blue')
lines(10^(pred$pred+2*pred$se),col='orange')
lines(10^(pred$pred-2*pred$se),col='orange')
```

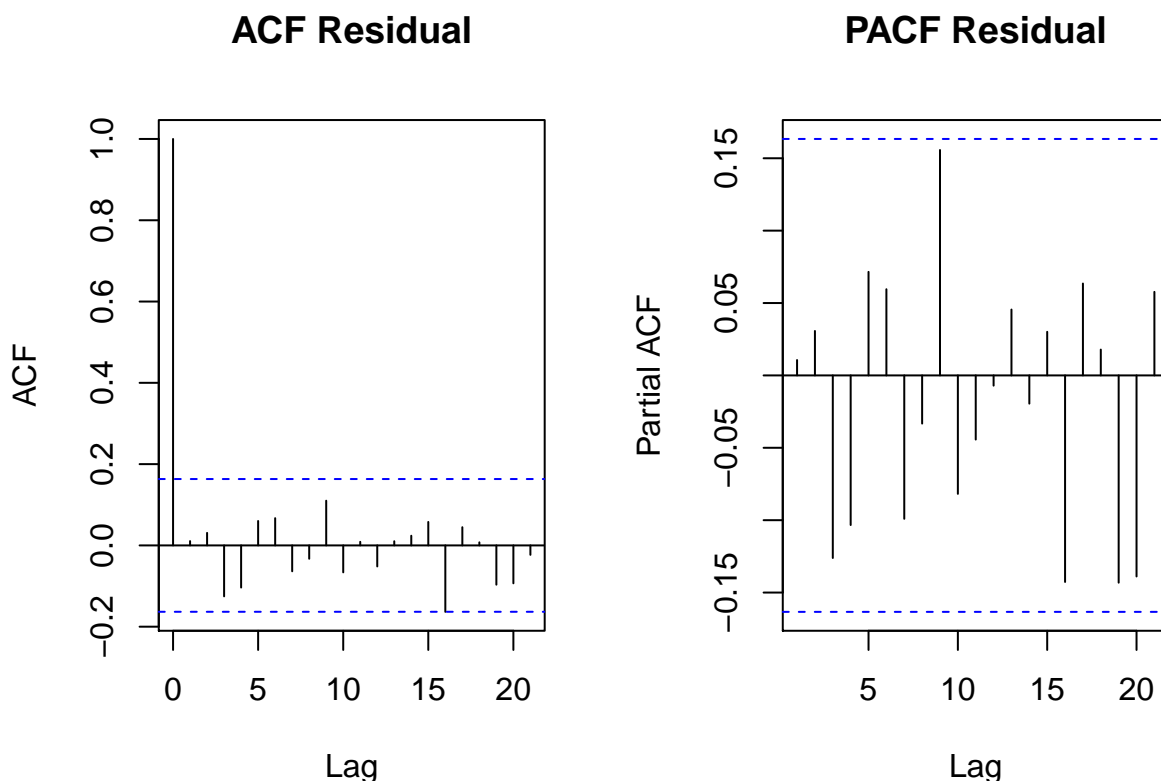


The above is the output with forecasted values of tractor sales in blue. Also, the range of expected error (i.e. 2 times standard deviation) is displayed with orange lines on either side of predicted blue line.

Assumptions while forecasting: Forecasts for a long period of 3 years is an ambitious task. The major assumption here is that the underlining patterns in the time series will continue to stay the same as predicted in the model. A short-term forecasting model, say a couple of business quarters or a year, is usually a good idea to forecast with reasonable accuracy. A long-term model like the one above needs to be evaluated on a regular interval of time (say 6 months). The idea is to incorporate the new information available with the passage of time in the model.

Plot ACF and PACF for residuals of ARIMA model to ensure no more information is left for extraction:

```
par(mfrow=c(1,2))
acf(ts(ARIMAFit$residuals),main='ACF Residual')
pacf(ts(ARIMAFit$residuals),main='PACF Residual')
```



Since there are no spikes outside the insignificant zone for both ACF and PACF plots we can conclude that residuals are random with no information or juice in them. Hence our ARIMA model is working good and predictions were successfully made.

Inference:

- Tractor sales seem to have both seasonality and trend.
- The sales are on average more during the months of July and August compared to other months by about 23% of annual trend.
- If the company has a high budget, we can confidently plan the production to be the forecasted sales $\pm 10\%$.

- If costs are high and company has a low budget, then we can plan the production to be 450 and be confident all tractors are going to be sold.
- Furthermore, a cost matrix and cost curve can be used to minimize the costs considering both cases.